

ATUL BHARDWAJ

(206) 609-4655 | atuloninternet@icloud.com | [linkedin.com/in/atul1000](https://www.linkedin.com/in/atul1000) | github.com/4uatul | atulbhardwaj.com

No immediate sponsorship required via. STEM OPT through mid-2029. Open to relocation.

Education

Seattle University

Master of Science, Computer Science (Data Science Specialization), GPA: 3.6

Washington, USA

Sep 2024 - Jun 2026

Maharshi Dayanand University

Bachelor of Technology, Computer Science Engineering, GPA: 3.5

Delhi, India

Aug 2020 - May 2024

Technical Skills

Programming Languages: Python, JavaScript, SQL, Bash

Web Development: React, Flask, PostgreSQL, REST APIs

Cloud & DevOps: AWS (3x Certified), Git, GitHub Actions, Docker, pytest, Code Review, Claude Code, GitHub Copilot

AWS Services: Lambda, API Gateway, S3, EC2, RDS, DynamoDB, VPC, IAM, Bedrock

AI/ML: Prompt Engineering, Multi-Agent Orchestration, RAG, Bedrock Knowledge Bases, AgentCore, Strands Agents SDK, MCP, Responsible AI

Data & Analytics: pandas, NumPy, Jupyter, Spark, AWS Glue, Athena, Matplotlib

Relevant Certifications

[AWS Certified Solutions Architect - Associate \(SAA-C03\)](#)

[AWS Certified Cloud Practitioner \(CLF-C02\)](#)

[AWS Certified AI Practitioner \(AIF-C01\)](#)

AWS AI & ML Scholars - Future AWS Agent Engineer Nanodegree, Udacity (in progress)

Professional Experience

Cloud Solutions Lead

AWS Cloud Club, Seattle University

Seattle, WA

Jan 2026 - Jun 2026

- Grew the AWS Cloud Club from **13 to 90+** members by building the technical programming from scratch. Guest speaker at the AWS Skills Center AI Exploration Day (**160 attendees**).
- Designed three tracks for AWS Hackathon 2026 (**Bedrock, Serverless, Cloud for Good**) and ran prep workshops. Secured \$10K in AWS credits through AWS VP Jeff Barr. **100+ participants, 20+ final projects**.
- Architected a **three-tier application** (CloudFront, EC2, RDS), **evaluated Lambda/DynamoDB against EC2/RDS**. Designed a **custom VPC** with subnet isolation and **application-tier security groups**. **\$0.90/day**.

Data and Operations Technician

Campus Services, Seattle University

Seattle, WA

Oct 2024 - Jun 2026

- Replaced 53 static SharePoint pages across 5 departments with a **fully serverless RAG application** (Bedrock, Lambda, API Gateway, S3) answering **20+ staff questions daily**. Tuned prompts and retrieval against a validation set of real questions; wrote the runbook and setup notes.
- Analyzed **100K+ rows** of parking and transit data in **Python and Tableau**, identified 4 **underused parking** and 3 **transit** plans; approved changes simplified user-facing options and saved ~5 hours per quarter.
- Rewrote the **Entrust batch SQL** to auto-sort thousands of ID cards at query time, cutting a manual step of **~30 minutes per batch**.

Data Science Consultant

Statistics Without Borders

Remote

Sep 2025 - Apr 2026

- Ran weekly syncs as technical lead across **7 teams (14 contributors, 14+ datasets)**, reviewing code and notebooks and enforcing consistency.
- Migrated 14 contributors from email workflows to **GitHub**. Authored onboarding docs and a pull-request review workflow that got **non-technical members committing in under a week**.

Hands-On Projects

AWS Service Advisor Agent (Bedrock AgentCore, Strands SDK, MCP) - 2nd Place, AWS Agentic AI GameDay

[Github](#)

- Built a **multi-agent system on Bedrock** where an orchestrator agent reasons through requirements and routes to specialist tools and a **cost analyzer agent**. Engineered the system prompts and tool schemas that drive routing.
- Deployed on **AgentCore's managed harness** with a custom MCP pricing server bridged through Gateway, **choosing Amazon Nova Micro** to cut per-query cost 90%. Validated agent behavior against scripted sample queries before each deploy.

YouTube Trending Analytics Data Lake (Amazon S3, AWS Glue, Lambda, Athena, PySpark, IAM, Python)

[Github](#)

- Built a serverless data lake on AWS to process **1,000+ YouTube trending videos across 5 regions** using S3 tiered storage, Glue for PySpark ETL, Lambda for JSON-to-Parquet conversion, and Athena for SQL.
- Cut storage by **75%** with Parquet and query cost **80%** via Hive-style partitioning. Automated the pipeline with S3-triggered Lambda and least-privilege IAM roles.